

# MOVIE RATING PREDICTION

This project involves building a machine learning model to predict movie ratings based on various features such as **genre, director, actors**, and other metadata. The project focuses on analyzing historical movie data, preprocessing it, engineering features, and applying regression techniques.

## Objective

- To build a machine learning model that predicts the **rating of a movie** based on its features.
- Analyze the relationship between features (genre, director, actors, etc.) and ratings given by **users** or **critics**.
- Provide insights into the factors influencing movie ratings.

## Data Dictionary

**Name:** The title of the movie.

**Year:** The year the movie was released.

**Genre:** The primary genre(s) of the movie (e.g., Action, Comedy, Drama). Multiple genres may be included.

**Rating:** The average rating of the movie given by critics or users (e.g., IMDb or Rotten Tomatoes rating).

**Votes:** The total number of votes or reviews that contributed to the movie's rating.

**Director:** The name of the director who directed the movie.

**Actor 1:** The name of the lead actor in the movie.

**Actor 2:** The name of the second lead actor or supporting actor.

**Actor 3:** The name of another key actor, often a supporting role.

## Task Steps

1. Setup Environment
2. Data Exploration
3. Data Cleaning
4. Feature Engineering
5. Split Data
6. Model Building
7. Model Evaluation
8. Insights and Visualization

## Python Packages

- **Pandas** (data manipulation)
- **Numpy** (numerical operations)
- **Matplotlib** (basic plotting)
- **Seaborn** (statistical visualizations)
- **scikit-learn** (for preprocessing, regression models, and evaluation)

## # 1. Loading the Dataset and Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_squared_log_error, r2_score
```

## # 2. Load the dataset

```
df = pd.read_csv('IMDB Movies India.csv', encoding='latin1')

# Dataset first look
print('First few rows:\n',df.head(5))
print('\nDataset Information:\n',df.info())
```

First few rows:

	Name	Year	...	Actor 2	Actor 3
0		NaN	...	Birbal	Rajendra Bhatia
1	#Gadhvi (He thought he was Gandhi)	(2019)	...	Vivek Ghamande	Arvind Jangid
2	#Homecoming	(2021)	...	Plabita Borthakur	Roy Angana
3	#Yaaram	(2019)	...	Ishita Raj	Siddhant Kapoor
4	...And Once Again	(2010)	...	Rituparna Sengupta	Antara Mali

[5 rows x 10 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 15509 entries, 0 to 15508

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	Name	15509 non-null	object
1	Year	14981 non-null	object
2	Duration	7240 non-null	object
3	Genre	13632 non-null	object

4 Rating 7919 non-null float64  
5 Votes 7920 non-null object  
6 Director 14984 non-null object  
7 Actor 1 13892 non-null object  
8 Actor 2 13125 non-null object  
9 Actor 3 12365 non-null object  
dtypes: float64(1), object(9)  
memory usage: 1.2+ MB

Dataset Information:

None

### # 3. Data Cleaning

```
# Initial DataFrame inspection
print("\nMissing Values Before Cleaning:\n",df.isnull().sum())

# Checking Duplicate Rows
print("\nDuplicate Rows Before Cleaning:\n",df.duplicated().sum())

# Drop rows with missing values
df.dropna(inplace=True)

# Recheck for missing values
print("\nMissing Values After Dropping:\n",df.isnull().sum())

# Drop duplicate rows
df.drop_duplicates(inplace=True)

# Final shape of the DataFrame
print("Final Shape of the DataFrame:", df.shape)
```

Missing Values Before Cleaning:

Name	0
Year	528
Duration	8269
Genre	1877
Rating	7590
Votes	7589
Director	525
Actor 1	1617
Actor 2	2384
Actor 3	3144

dtype: int64

Duplicate Rows Before Cleaning:

6

Missing Values After Dropping:

Name 0

Year 0

Duration 0

Genre 0

Rating 0

Votes 0

Director 0

Actor 1 0

Actor 2 0

Actor 3 0

dtype: int64

Final Shape of the DataFrame: (5659, 10)

#### # 4. Data Pre-processing

```
# replacing the brackets from year column
df['Year'] = df['Year'].str.replace(r'[\(\)]', '', regex=True).astype(int)

# Remove the min word from 'Duration' column and convert all values to numeric
df['Duration'] = pd.to_numeric(df['Duration'].str.replace(' min', ''))

# Splitting the genre by, to keep only unique genres and replacing the null values with
mode
df['Genre'] = df['Genre'].str.split(', ')
df = df.explode('Genre')
df['Genre'].fillna(df['Genre'].mode()[0], inplace=True)

# Convert 'Votes' to numeric and replace the to keep only numeric part
df['Votes'] = pd.to_numeric(df['Votes'].str.replace(', ', ''))

# Checking the dataset is there any null values present and data types of the feature
present
print('\nDataset Information:\n', df.info())
```

FutureWarning: A value is trying to be set on a copy of a DataFrame

or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or  
df[col] = df[col].method(value) instead, to perform the operation inplace on the original object

```
df['Genre'].fillna(df['Genre'].mode()[0], inplace=True)
```

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 11979 entries, 1 to 15508

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

--- -----

0	Name	11979 non-null	object
---	------	----------------	--------

1	Year	11979 non-null	int64
---	------	----------------	-------

2	Duration	11979 non-null	int64
---	----------	----------------	-------

3	Genre	11979 non-null	object
---	-------	----------------	--------

4	Rating	11979 non-null	float64
---	--------	----------------	---------

5	Votes	11979 non-null	int64
---	-------	----------------	-------

6	Director	11979 non-null	object
---	----------	----------------	--------

7	Actor 1	11979 non-null	object
---	---------	----------------	--------

8	Actor 2	11979 non-null	object
---	---------	----------------	--------

9	Actor 3	11979 non-null	object
---	---------	----------------	--------

dtypes: float64(1), int64(3), object(6)

memory usage: 1.0+ MB


Dataset Information:

None

## # 5. Data Visualizing

### # Creating histogram plot

```
# Here we have created a histogram over the years in the data
plt.hist(df['Year'], bins=30, density=True, alpha=0.7, color='blue')
plt.title('Distribution of Movies by Year')
plt.xlabel('Year')
plt.ylabel('Probability Density')
plt.show()
```



```
# Group data by Year and calculate the average rating
avg_rating_by_year = df.groupby(['Year', 'Genre'])['Rating'].mean().reset_index()
print(avg_rating_by_year.head(10))

# Get the top 10 genres by average rating
top_genres =
avg_rating_by_year.groupby('Genre')['Rating'].mean().sort_values(ascending=False).head(10).
index
print(top_genres)

# Filter the data to include only the top 3 genres
average_rating_by_year = avg_rating_by_year[avg_rating_by_year['Genre'].isin(top_genres)]
```

```
print(average_rating_by_year.head(10))
```


	Year	Genre	Rating
0	1931	Drama	5.75
1	1931	Fantasy	6.20
2	1932	Musical	6.00
3	1932	Romance	6.00
4	1933	Drama	6.20
5	1933	Romance	6.20
6	1934	Adventure	2.70
7	1934	Drama	8.50
8	1934	Fantasy	2.70
9	1935	Action	4.50

```
Index(['News', 'Documentary', 'Biography', 'History', 'Musical', 'Drama',  
      'Family', 'Sport', 'Music', 'Comedy'],  
      dtype='object', name='Genre')
```

	Year	Genre	Rating
0	1931	Drama	5.750000
2	1932	Musical	6.000000
4	1933	Drama	6.200000
7	1934	Drama	8.500000
11	1935	Drama	6.633333
13	1936	Drama	6.350000
14	1936	Musical	6.500000
17	1937	Biography	6.700000
18	1937	Drama	6.525000
19	1937	Family	5.900000


### # Creating line plot

```
plt.figure(figsize=(12, 8))
sns.lineplot(data=average_rating_by_year, x='Year', y='Rating', hue='Genre',
marker='o')
plt.title('Average Rating by Genre Over Time')
plt.xlabel('Year')
plt.ylabel('Average Rating')
plt.show()
```




```
# This histogram shows the distribution of ratings and its probability density
plt.hist(df['Rating'], bins=30, density=True, alpha=0.7, color='blue')
plt.title('Distribution of Ratings')
plt.xlabel('Rating')
plt.ylabel('Probability Density')
plt.show()
```





### # Creating a scatter plot

```
# This scatter plot shows the relationship between the duration of the movie and its
rating
plt.figure(figsize=(12, 8))
sns.scatterplot(data=df, x='Duration', y='Rating', hue='Genre')
plt.title('Duration vs Rating')
plt.xlabel('Duration')
plt.ylabel('Rating')
plt.show()
```



## # 6.Feature Engineering

```
# Drop the Name column because it doesn't impact the outcome
df.drop('Name', axis=1, inplace=True)

# Grouping the columns with their average ratings and then creating a new feature
genre_mean_rating = df.groupby('Genre')['Rating'].transform('mean')
df['Genre_mean_rating'] = genre_mean_rating

director_mean_rating = df.groupby('Director')['Rating'].transform('mean')
df['Director_encoded'] = director_mean_rating

actor1_mean_rating = df.groupby('Actor 1')['Rating'].transform('mean')
df['Actor1_encoded'] = actor1_mean_rating

actor2_mean_rating = df.groupby('Actor 2')['Rating'].transform('mean')
df['Actor2_encoded'] = actor2_mean_rating

actor3_mean_rating = df.groupby('Actor 3')['Rating'].transform('mean')
df['Actor3_encoded'] = actor3_mean_rating

# Display the updated DataFrame
print(df.head())
```

	Year	Duration	Genre	Rating	Votes	...	Genre_mean_rating	Director_encoded	Actor1_encoded	Actor2_encoded	Actor3_encoded
1	2019	109	Drama	7.0	8 ...	6.056744	7.000000	6.850000	7.00	7.00	
3	2019	110	Comedy	4.4	35 ...	5.751042	4.400000	5.250000	4.40	4.46	
3	2019	110	Romance	4.4	35 ...	5.811087	4.400000	5.250000	4.40	4.46	
5	1997	147	Comedy	4.7	827 ...	5.751042	5.335135	4.793617	5.73	5.93	
5	1997	147	Drama	4.7	827 ...	6.056744	5.335135	4.793617	5.73	5.93	

[5 rows x 14 columns]

## # 7. Split Data

```
# Keeping the predictor and target variables
X = df[['Year', 'Votes', 'Duration', 'Genre_mean_rating', 'Director_encoded',
'Actor1_encoded', 'Actor2_encoded', 'Actor3_encoded']]
y = df['Rating']

# Splitting the data into training and testing data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
print(f"Train shapes: X_train: {X_train.shape}, y_train: {y_train.shape}")
```

```
print(f"Test shapes: X_test: {X_test.shape}, y_test: {y_test.shape}")
```

Train shapes: X\_train: (9583, 8), y\_train: (9583,)

Test shapes: X\_test: (2396, 8), y\_test: (2396,)

## # 8. Model Building

```
# Initializing and training the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Making predictions
y_pred = model.predict(X_test)

# Calculating evaluation metrics
print('Accuracy (R2 Score):', model.score(X_test, y_test))
print('Mean Squared Error:', mean_squared_error(y_test, y_pred))

# Handling potential errors for negative predictions in mean_squared_log_error
y_pred = [max(val, 0) for val in y_pred]
print('Mean Squared Log Error:', mean_squared_log_error(y_test, y_pred))

print('R2 Score:', r2_score(y_test, y_pred))
```

Accuracy (R2 Score): 0.7641133663863862

Mean Squared Error: 0.4465441653985704

Mean Squared Log Error: 0.012458877846073142

R2 Score: 0.7641133663863862

## # Model Testing

```
print(X.head())
print(y.head())

# For testing , we create a new dataframe with values close to the any of our existing data
to evaluate
data = {'Year': [2019], 'Votes': [36], 'Duration': [111], 'Genre_mean_rating': [5.8],
'Director_encoded' : [4.5], 'Actor1_encoded' : [5.3], 'Actor2_encoded' : [4.5],
'Actor3_encoded' : [4.5]}
trail = pd.DataFrame(data)

# Predict the movie rating by entered data
rating = model.predict(trail)
```

```
# Display the predicted rating
print('Predicted Rating:', rating[0])
```

	Year	Votes	Duration	Genre_mean_rating	Director_encoded	Actor1_encoded	Actor2_encoded	Actor3_encoded
1	2019	8	109	6.056744	7.000000	6.850000	7.00	7.00
3	2019	35	110	5.751042	4.400000	5.250000	4.40	4.46
3	2019	35	110	5.811087	4.400000	5.250000	4.40	4.46
5	1997	827	147	5.751042	5.335135	4.793617	5.73	5.93
5	1997	827	147	6.056744	5.335135	4.793617	5.73	5.93
1							7.0	
3							4.4	
3							4.4	
5							4.7	
5							4.7	

Name: Rating, dtype: float64

Predicted Rating: 4.2074589621343295

## Conclusion:

The movie rating prediction program provides a comprehensive approach to analyze, clean, and model movie data to predict ratings based on various features like year, votes, duration, genre, director, and actors.

## Insights:

- The program highlights the importance of feature engineering and cleaning in building effective predictive models.
- Linear regression works well for predicting continuous variables like movie ratings, but its performance could be further enhanced using more complex models like Random Forest or Gradient Boosting if required.